

CS4103-DS: Security

Jan de Muijnck-Hughes

11 April 2016, 12 April 2016, 19 April 2016

Section 1

Overview

Aims & Objectives

- Gain an understanding of salient issues surrounding **Security** and Distributed Systems.
- Understand the issues associated with **authorisation** within a Distributed System, and ways in which it can be addressed.
- Understand issues associated with **authentication**, and how cryptographic techniques can be used to provide authentication mechanisms.

Reading

- Andrew Tanenbaum et al. *Distributed Systems: Principles and Paradigms*. English. 3rd ed. Pearson Higher Education, 2013, p. 633. ISBN: 1292025522, Chp. 9:§9.1-2, §9.2.1-2&4 §9.3.1, §9.4.1&3. §9.5
- George Coulouris et al. *Distributed Systems: Concepts and Designs*. English. 5th ed. Pearson Higher Education, 2011, p. 927. ISBN: 0273760599, Chp. 11:§11.1, §11.6.1&2
- Yu Zhou et al. 'Policy Enforcement Pattern'. In: *PLoP 2002*. 2002

Security as Risk Management

Doing Security \equiv Risk Management

- Asset identification
- Risk identification
 - Identifying an asset's vulnerabilities
 - Identifying relevant threats
- Risk analysis
- Risk treatment

Threat Manifestation aka Risk

$$\text{Risk} \leftarrow \text{Threat} + \text{Vulnerability} = \text{Success}$$

Threat

- Circumstances that have potential to cause loss or harm to the asset
- Threats can be:
 - **accidental**
 - **deliberate**
 - **environmental**

Vulnerability

- Weakness that can be exploited within a system
- Vulnerabilities can be:
 - **accidental**
 - **deliberate**
 - **environmental**

ISO Threat Types

- **Physical damage**
 - fire, water, dust
- **Natural events**
 - weather, volcanic activity
- **Loss of essential services**
 - loss of power
- **Disturbance due to radiation**
 - electromagnetic, thermal
- **Compromise of information**
 - Eavesdropping, Remote Spying
- **Technical failures**
 - equipment or software malfunction
- **Unauthorised Actions**
 - illegal processing of data, using pirated software
- **Compromise of functions**
 - Abuse of rights, Denial of Actions

Where can Vulnerabilities Occur?

- **Hardware**
 - environmental damage, wear and tear
- **Software**
 - well-known flaws, insufficient testing
- **Network**
 - single point of failure, unprotected comm lines
- **Personnel**
 - lack of personnel, insufficient training
- **Site**
 - located in flood plain, unstable power grid
- **Organisational**
 - lack of continuity plans, lack of email usage policy

Security Policies & Mechanisms

Policies

Describes the actions that an 'entity' are permitted to do, and not to. Essentially, security requirements: Confidentiality, Integrity, Availability...

- **Examples**

- 'Only Jan & DoT can see the exam.'
- 'STAFFRESS is only accessible by Staff members.'

Mechanisms

Technology or procedure employed to enforce the policy.

- **Examples**

- Authentication, Authorisation, Encryption, & Auditing.

Intermezzo: Scope

- Security is a comprehensive and extensive subject area.
- Our interest for this lecture is:

Security of Distributed Systems

and

Distributed Systems for Security.

- We won't cover other security topics.

Section 2

Security & Distributed Systems

Security & Distributed Systems

How can security policies be defined and implemented over distributed resources and using what mechanisms?

Core Issues concern Identity & Access Management

1 Data Security:

- How to secure data **at-rest**?
- How to secure data **in-flight**?

2 Identity Management:

- Definition and management of identities.

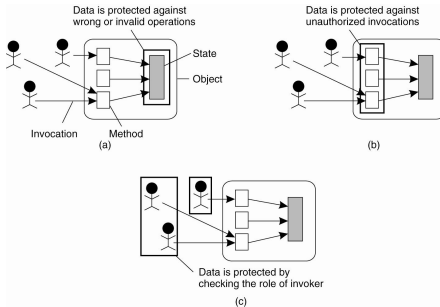
3 Authentication:

- Authentication in distributed setting.

4 Authorisation:

- Define and enforce authorisation policies.
- Authorisation in distributed setting.

Focus of Control




Tanenbaum et al. [TS13]

Where to focus protection?

- 1 **Model** Protection against invalid operations
- 2 **View** Protection against unauthorised invocations
- 3 **Controller** Protection against unauthorised users


Layering of Security Mechanisms



Tanenbaum et al. [TS13]

- Protect all the layers.
- 'Transport Layer Security is solved'

Distribution of Security Mechanisms



Tanenbaum et al. [TS13]

- Organisational & Administrative Heterogeneity.
- Trusted Computing Base
- Simplicity

Securing Data at Rest

Data at rest is data that doesn't 'move'.

- Can solve using cryptography...
 - **Secrecy with Signature with Appendix** using **KEM/DEM**
 - Public Key Encryption, Symmetric Encryption, Digital Signatures, Hash functions...
- **What standards and parameters to use?**

Disk, File-System	Symmetric	AES, SkipJack, Blowfish
User files, app data	KEM/DEM	RSA, DSA, ECC, ECDSA
- **Key Management!?**
 - Public Key Infrastructure: Centralised, Decentralised
- **Expressiveness of Encryption**
 - Perfect Forward Secrecy
 - Classical Schemes provide **1-2-1** Encryption.
 - Need additional mechanisms to manage permissions.

Securing Data in Flight

Data in flight is data being moved from one domain to another.

- Can solve using cryptography to construct **secure channels**
 - Send messages securely between two points: End-2-End Encryption.
- What standards and parameters to use?
 - Network Layer has IPSec
 - Transport Layer has TLS
 - Application layer has: Signal, Cryptocat, OTR...
- Key Management!?
 - Public Key Infrastructure: Centralised, Decentralised

Section 3


Authentication

Authentication

Given two entities Alice and Bob, how can Bob authenticate with Alice such that Alice knows that Bob is really who he says he is.

- Typically 'Two phased' protocols
 - 1 **Enrollment**: Establish credentials.
 - 2 **Challenge & Response**: Check validity of credentials.
- Utilise lots of cryptography.
- **Styles**
 - Direct or Brokered
- **Examples**
 - Network: HIP, IPSec, ILNP
 - Transport: MS-CHAP, EAP
 - Application: RADIUS, DIAMETER
 - User: SAML, OpenID, .Net Passport, KERBEROS, Shibboleth, OpenAthens

Secure Channel



- Use **Secrecy with Signature with Appendix** using **KEM/DEM**
 - Session Keys for each conversation.
 - Public Key Infrastructure to get public keys.
- How to exchange session keys?
 - Diffie-Hellman Key Exchange, Station-To-Station, Needham-Schroeder-Lowe
- How to authenticate Bob?

Kerberos

Authentication protocol using **tickets** to allow nodes to authenticate over an untrusted network.

- Developed by MIT.
- Requires a **Trusted-Third Party**
 - Authentication Service
 - Ticket Granting Service
- Mutual Authentication
- 'Dated'

Intermezzo: Crypto Notation

Key Notation

Symmetric Key K_{AB}

Public Key $Enc_{pub}(Bob)$

Private Key $Dec_{priv}(Bob)$

Signing Key $Enc_{priv}(Alice)$

Verifying Key $Dec_{pub}(Alice)$

Operations

Encrypt $Encrypt(\dots)$

Decrypt $Decrypt(\dots)$

Sign $Sign(\dots)$

Verify $Verify(\dots)$

Misc

Ctxt Sym $\{M\}_{K_{Bob}}$

Hash $\#(msg)$

Concatenate $A || B$

Ctxt ASym $\{|M|\}_{Enc(Bob)}$

Send A to B $A \rightarrow B : msg$

Assignment $H_{msg} \leftarrow \#(msg)$

Authentication Protocol

Simplified Kerberos protocol to talk to Bob.

- **Sign into Service**
 - Session Key $K_{A,AS}$ Established
 - $Alice \rightarrow AS : ID(A)$
 - AS generates
 - ticket with TTL: $\mathcal{T}_{ttl} \leftarrow \{ID(A) \parallel K_{A,TGS}\}_{K_{AS,TGS}}$
 - Session Key $K_{A,TGS}$
 - $AS \rightarrow Alice : \{K_{A,TGS} \parallel \mathcal{T}_{ttl}\}_{K_{A,AS}}$
- **Request Ticket to Talk to Bob**
 - Session Key $K_{A,TGS}$ Established
 - Timestamp t .
 - $A \rightarrow TGS : \mathcal{T}_{ttl} \parallel ID(B) \parallel \{t\}_{K_{A,TGS}}$
 - TGS Generates Session Key $K_{A,B}$ and obtains $K_{B,TGS}$.
 - $TGS \rightarrow A : \{ID(B) \parallel K_{A,B}\}_{K_{A,TGS}} \parallel \{ID(A) \parallel K_{A,B}\}_{K_{B,TGS}}$
- **Ask Bob To Talk**
 - $A \rightarrow B : \{ID(A) \parallel K_{A,B}\}_{K_{B,TGS}} \parallel \{t\}_{K_{A,B}}$
 - $B \rightarrow A : \{t+1\}_{K_{A,B}}$

Kerberos cont. . .

- **Authentication Protocol**
 - Based on Needham-Schoeder-Lowe
- **‘Single-Sign-On’**
 - By authenticating with the AS get timed access (24hrs) to system.
 - Ticket used to request access to other services i.e. other bobs.
 - Combine with Authorisation services
- **‘Simplified’**
 - Introduce Public Key variants
 - Don't see sending $\{\{ID(A) \parallel K_{A,B}\}_{K_{B,TGS}} \parallel \{t\}_{K_{A,B}}\}_{Enc(B)}$

Advantages & Disadvantages

- Advantages
 - Authentication in a Distributed System
 - Single-Sign-On
- Disadvantages
 - Single Point of Failure
 - Not federated
 - 'Dated'
 - Not a cool protocol. . .

Section 4

Authorisation

Authorisation/Access Control

Granting access rights to a subject for resources in various environments, and ensuring that a subject has the correct permissions to access a particular resource in an particular environment.

- **Access Control Models**

- Access Control Matrix
- Access Control Lists, Capabilities
- Role-Based Access Control, Attribute-Based Access Control, Policy-Based Access Control

- **Implementations**

- POSIX, Capsicum, XACML, SAML, Kerberos, Shibboleth, OpenID, OAuth, Facebook Connect

General Model



- **Subjects** are: nodes, processes, users. . .
- **Objects** are: files, data, databases, services. . .
- **Permissions** are actions on objects: Read, Write, Execute. . .
 - Permissions can also be time dependent.
- **Monitor** is access control mechanism to enforce permissions.
- **Schema** is a description of an instance of an access control model for a particular scenario.

Access Control Matrix


Matrix where rows denote subjects, columns denote objects, and cells the permissions that the subject has on an object.

	Slides	Exam	STAFFRES
Jan	RWX	RWX	RWX
DoT	R-X	RWX	RWX
Bob	R-X	—	—

- Common way to envisage access control.
- Monitor 'just' performs matrix look up.
- If Subject s or Object o not in Matrix M then failure.
- Unwieldy for large models

Access Control Lists


Each object carries a description of the subjects and their permissions.



- Classic Approach found in most OS.
 - Column Spans of a matrix.
- Monitor/Object **needs to know** who can do what
- 'Centralised' Approach.


Capabilities

Each subject carries a description of the objects they can access and their assigned permissions.



- Row Span of a Matrix
- Each client is given restricted list of abilities on objects.
- Monitor checks if capability can be applied.
- 'Decentralised' Approach.


Role-Based Access Control



Application Layer Model

- Each subject has one or more (Hierarchic) Roles.
- Object has permissions based on roles.
- Monitor checks if Subject's Role allows access to Object.
- 'Decentralised' Approach.

Attribute-Based Access Control



Application Layer Model

- Attributes used to describe: **Subjects, Objects, & the Environment.**
- Policies are **Boolean Formula** over attributes.
- Monitor grants access based on policy satisfaction.
- 'Decentralised' Approach.

ABAC Example: TCPLog Access

Attributes

Subject Group, Roles, Clearance Level...

Object TCP Header Information, Ownership...

Environment Locale, Time, Date...

Access Policy

$$\begin{aligned} \text{Policy}(s, o, e) \leftarrow & \text{Group}(s) \equiv \text{GCHQ} \\ & \wedge \text{Level}(s) \geq \text{SECRET} \\ & \wedge (\text{srcPort}(o) \equiv 80 \vee \text{srcPort}(o) \equiv 8080) \\ & \wedge \text{srcAddr}(o) \equiv 123.456.789 \\ & \wedge \text{CurrentDate}(e) \leq 20160527 \\ & \wedge \text{CurrentDate}(e) \geq 20150927 \end{aligned}$$

XACML: eXtensible A/C Mark-up Language

Declarative access control policy language and processing model using XML to encode and evaluate policies.


- OASIS Standard [Ris13].
- Policy Language is based on ABAC.
- 'Policy Enforcement Points'
 - Separates decision from enforcement from definition.
 - Distributed components that work together.
- Designed for **Service Oriented Architectures**

Policy Enforcement Points

General architectural model to describe a scalable distributed authorisation framework.

- Described as a **Design Pattern** in Zhou et al. [ZZP02].
- Generalisation of AAA Framework [Vol+00]
- **Key Features**
 - Distributed components that work together.
 - Separates decision from enforcement from definition.
 - Policies are made on demand, or pre-made.
 - Policies are taken from ABAC

PEP Architecture



PEP Architecture

Taken from Rissanen [Ris13]

- **PDP** Policy Decision Point
 - The system entity that evaluates applicable policy and renders an authorization decision.
- **PEP** Policy Enforcement Point
 - The system entity that performs access control, by making decision requests and enforcing authorization decisions.
- **PIP** Policy Information Point
 - The system entity that acts as a source of attribute values.
- **PAP** Policy Administration Point
 - The system entity that creates a policy or policy set.

Section 5

Summary

Summary

- Security is **hard**; Security is a socio-technical problem.
- Four 'core' security issues for Distributed systems:
 - **Data Security**: In Flight, At Rest.
 - **Identity Management**: Describing and managing entities.
 - **Authentication**: Verify entities identity.
 - **Authorisation**: Verify their permissions.
- Establishing **Secure Channels** often requires brokered authentication.
- Access Control Models help manage permissions at OS and Application Level.
- **Policy Enforcement Points** design pattern to provide distributed access control.

Section 6

Crypto Basics

Why use Cryptography?

Cryptography can be used to provide **mathematical** guarantees towards:

- **Confidentiality**
 - Public Key Encryption i.e. RSA, ElGamal, ECC
 - Block Ciphers i.e. Blowfish, TripleDES, Skipjack, AES
 - Stream Ciphers i.e. RC4
- **Integrity**
 - Cryptographic Hash function i.e. MD-family, SHA-family
 - Message Authentication Codes
- **Authenticity & Non-Repudiation**
 - Digital Signatures i.e. DSS, (EC)DSA

Cryptographic Hash Functions/Message Digests

Definition

Function to compute a unique (random) signature for some data:

$$\# : \{0, 1\}^n \rightarrow^R \{0, 1\}^n$$

- **Provides guarantees towards:** Data Integrity.
- **Properties**
 - **Pre-image resistance:** Given $\#(m)$, hard to find m .
 - **Second Pre-image resistance:** Given m_1 , hard to find m_2 such that $m_1 \neq m_2$ & $\#(m_1) == \#(m_2)$
 - **Collision Resistance:** Hard to find m_1, m_2 such that $\#(m_1) == \#(m_2)$
- **Implementations**
 - MD-Family, SHA-Family

Block Ciphers: Symmetric Cryptography

Definition

Set of functions to encrypt data.

$$C \leftarrow \text{Encrypt}(M, K_M)$$

$$M \leftarrow \text{Decrypt}(C, K_M)$$

- **Provides guarantees towards:** Confidentiality
- **Properties**
 - Same key used to encrypt and decrypt.
 - Implementations are very efficient for large messages.
- **Implementations**
 - Blowfish, TripleDES, Skipjack, AES

Asymmetric Ciphers

Definition

Set of functions to encrypt data.

$$\begin{aligned}(\text{Enc}(\text{Alice}), \text{Dec}(\text{Alice})) &\leftarrow \text{KeyGen}(\lambda) \\ C &\leftarrow \text{Encrypt}(M, \text{Enc}(\text{Alice})) \\ M &\leftarrow \text{Decrypt}(C, \text{Dec}(\text{Alice}))\end{aligned}$$

- **Provides guarantees towards:** Confidentiality
- **Properties**
 - Use of Key Pairs.
 - One key used to encrypt, the other decrypt.
 - Two modes of use: Encrypting & Signing
 - Very inefficient on large data.
- **Implementations**
 - DSA, (EC)DSA, ECC, RSA, ElGamal

Cryptographic Workflows

Ways in which crypto primitives can be combined/used to provide one or more security guarantees.

- **Information Secrecy**
- **Efficient Information Secrecy**
- **Sender Authentication**
- **Secrecy with Authentication**
- **Secrecy with Signature**
- **Secrecy with Integrity**
- **Signature with Appendix**
- **Secrecy with Signature with Appendix**

Public Key Cryptography

Asymmetric Crypto can be used to provide:

- **Confidentiality**

Key Generation

$$(\text{Enc}_{\text{pub}}(\text{Bob}), \text{Dec}_{\text{priv}}(\text{Bob})) \leftarrow \text{KeyGen}(\lambda)$$

Alice

- 1 $C \leftarrow \text{Encrypt}(M, \text{Enc}_{\text{pub}}(\text{Bob}))$
- 2 $\text{Alice} \rightarrow \text{Bob} : C$

Bob

- 1 $C' \leftarrow C$
- 2 $M' \leftarrow \text{Decrypt}(C', \text{Dec}_{\text{priv}}(\text{Bob}))$

Digital Signatures

Asymmetric Crypto and **Message Digests** can be used to provide:

- Authenticity of message origin
- Non-Repudiation of message origin
- Message Integrity

Key Generation

$$(\text{Enc}_{\text{priv}}(\text{Alice}), \text{Dec}_{\text{pub}}(\text{Alice})) \leftarrow \text{KeyGen}(\lambda)$$

Alice

- 1 $H \leftarrow \#(M)$
- 2 $S \leftarrow \text{Sign}(H, \text{Enc}_{\text{priv}}(\text{Alice}))$
- 3 $\text{Alice} \rightarrow \text{Bob} : (M \parallel S)$

Bob

- 1 $(M', S') \leftarrow (M \parallel S)$
- 2 $H' \leftarrow \text{Verify}(S', \text{Dec}_{\text{pub}}(\text{Alice}))$
- 3 Accept iff $\#(M') \equiv H'$

Public Key Encryption & Digital Signatures

Combining the previous primitives provides the following:

- Message Confidentiality and Integrity
- Authenticity of message origin
- Non-Repudiation of message origin

Alice

- 1 $H \leftarrow \#(M)$
- 2 $S \leftarrow \text{Sign}(H, \text{Enc}_{\text{priv}}(\text{Alice}))$
- 3 $C \leftarrow \text{Encrypt}(M, \text{Enc}_{\text{pub}}(\text{Bob}))$
- 4 $\text{Alice} \rightarrow \text{Bob} : (C \parallel S)$

Bob

- 1 $(C', S') \leftarrow (C \parallel S)$
- 2 $M' \leftarrow \text{Decrypt}(C', \text{Dec}_{\text{priv}}(\text{Bob}))$
- 3 $H' \leftarrow \text{Verify}(S', \text{Dec}_{\text{pub}}(\text{Alice}))$
- 4 Accept iff $\#(M') \equiv H'$

KEM/DEM

Improve encryption **efficiency** through a hybrid encryption scheme:

- **Symmetric** Encryption to encrypt **data**; and
- **Asymmetric** Encryption to encrypt **symmetric key**.

AKA **Key Encapsulation/Data Encapsulation** Mechanism

Alice

- 1 $C_M \leftarrow \text{Encrypt}(M, K_{\text{Random}})$
- 2 $C_K \leftarrow \text{Encrypt}(K_R, \text{Enc}_{\text{pub}}(\text{Bob}))$
- 3 $\text{Alice} \rightarrow \text{Bob} : (C_K \parallel C_M)$

Bob

- 1 $(C'_K, C'_M) \leftarrow (C_K \parallel C_M)$
- 2 $K'_R \leftarrow \text{Decrypt}(C'_K, \text{Dec}_{\text{priv}}(\text{Bob}))$
- 3 $M' \leftarrow \text{Decrypt}(C'_M, K'_R)$

Some Cryptographic Algorithms

- **RSA Systems**

- Security lies in the **hardness** of factorising large numbers.
- **Examples:** RSA Encryption, RSA Digital Signatures

- **Discrete Logarithm Systems**

- Security lies in the **hardness** of taking discrete logarithms over **finite fields**.
 - Key Exchange i.e. Diffie-Hellman Key Exchange
 - Digital Signature Algorithm
 - Discrete Logarithm Integrated Encryption Scheme

- **Note:** Many variants of DL schemes in different settings e.g. Elliptic Curves.

This list is far from complete. . .